

```
In [1]: from google.colab import files
uploaded = files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving salary.csv to salary.csv

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.preprocessing import LabelEncoder
```

/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.

```
import pandas.util.testing as tm
```

```
In [3]: df = pd.read_csv("salary.csv")
d = pd.read_csv("salary.csv")
```

In [4]: df

Out[4]:

	Unnamed: 0	age	Workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	nat
0	0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United States
1	1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United States
2	2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United States
3	3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United States
4	4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	United States
...
1027	1114	34	Private	290763	HS-grad	9	Divorced	Handlers-cleaners	Own-child	White	Female	0	0	40	United States
1028	1116	36	Private	51100	Some-college	10	Married-civ-spouse	Craft-repair	Husband	White	Male	0	0	40	United States
1029	1117	41	Private	227644	HS-grad	9	Married-civ-spouse	Transport-moving	Husband	White	Male	0	0	50	United States
1030	1118	58	Local-gov	205267	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	White	Female	0	0	40	United States
1031	1119	53	Private	288020	Bachelors	13	Married-civ-spouse	Prof-specialty	Husband	Asian-Pac-Islander	Male	0	0	40	United States


1032 rows × 16 columns

```
In [5]: df.drop(["Unnamed: 0"], axis=1,inplace = True)
```


```
In [6]: df["Income"].value_counts()
```

```
Out[6]: 0    778  
        1    254  
        Name: Income, dtype: int64
```

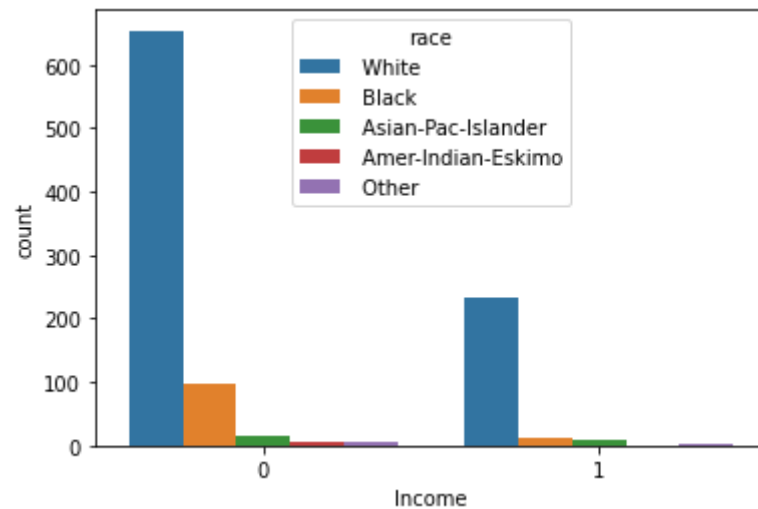
```
In [7]: sns.countplot(x='Income',data=df, palette='hls')  
plt.show()
```



```
In [8]: plt.figure()  
sns.countplot(data=df,x="Income",hue="sex")  
plt.show()
```



```
In [9]: plt.figure()  
sns.countplot(data=df,x="Income",hue="race")  
plt.show()
```



#To Check Missing Value:-

```
In [10]: df.isnull().sum()
```

```
Out[10]: age                0  
workclass              0  
fnlwgt                 0  
education              0  
education-num          0  
marital-status         0  
occupation             0  
relationship           0  
race                   0  
sex                    0  
capital-gain           0  
capital-loss           0  
hours-per-week         0  
native-country         0  
Income                 0  
dtype: int64
```

Data Type :

```
In [37]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1032 entries, 0 to 1031
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Workclass              1032 non-null   int64
1   education              1032 non-null   int64
2   marital-status         1032 non-null   int64
3   occupation             1032 non-null   int64
4   relationship           1032 non-null   int64
5   race                   1032 non-null   int64
6   sex                    1032 non-null   int64
7   native-country         1032 non-null   int64
8   age                    1032 non-null   int64
9   fnlwgt                 1032 non-null   int64
10  education-num          1032 non-null   int64
11  capital-gain            1032 non-null   int64
12  capital-loss            1032 non-null   int64
13  hours-per-week         1032 non-null   int64
14  Income                 1032 non-null   int64
dtypes: int64(15)
memory usage: 121.1 KB
```

```
In [12]: df_num = df.select_dtypes(["float64", "int64"])
```

```
In [13]: df_num
```

```
Out[13]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	Income
0	39	77516	13	2174	0	40	0
1	50	83311	13	0	0	13	0
2	38	215646	9	0	0	40	0
3	53	234721	7	0	0	40	0
4	28	338409	13	0	0	40	0
...
1027	34	290763	9	0	0	40	0
1028	36	51100	10	0	0	40	0
1029	41	227644	9	0	0	50	0
1030	58	205267	13	0	0	40	1
1031	53	288020	13	0	0	40	0

1032 rows × 7 columns

```
In [14]: df_cat = df.select_dtypes(object)
```



```
In [15]: df_cat
```

```
Out[15]:
```

	Workclass	education	marital-status	occupation	relationship	race	sex	native-country
0	State-gov	Bachelors	Never-married	Adm-clerical	Not-in-family	White	Male	United-States
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States
2	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States
3	Private	11th	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	United-States
4	Private	Bachelors	Married-civ-spouse	Prof-specialty	Wife	Black	Female	Cuba
...
1027	Private	HS-grad	Divorced	Handlers-cleaners	Own-child	White	Female	United-States
1028	Private	Some-college	Married-civ-spouse	Craft-repair	Husband	White	Male	United-States
1029	Private	HS-grad	Married-civ-spouse	Transport-moving	Husband	White	Male	United-States
1030	Local-gov	Bachelors	Married-civ-spouse	Prof-specialty	Wife	White	Female	United-States
1031	Private	Bachelors	Married-civ-spouse	Prof-specialty	Husband	Asian-Pac-Islander	Male	Japan

1032 rows × 8 columns

```
In [16]: le = LabelEncoder()
```

```
In [17]: for col in df_cat:
          le = LabelEncoder()
          df_cat[col] = le.fit_transform(df_cat[col])
```

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

This is separate from the ipykernel package so we can avoid doing imports until

```
In [18]: df_cat
```

```
Out[18]:
```

	Workclass	education	marital-status	occupation	relationship	race	sex	native-country
0	5	9	4	0	1	4	1	28
1	4	9	2	3	0	4	1	28
2	2	11	0	5	1	4	1	28
3	2	1	2	5	0	2	1	28
4	2	9	2	9	5	2	0	4
...
1027	2	11	0	5	3	4	0	28
1028	2	15	2	2	0	4	1	28
1029	2	11	2	13	0	4	1	28
1030	1	9	2	9	5	4	0	28
1031	2	9	2	9	0	1	1	18

1032 rows × 8 columns

```
In [19]: df = pd.concat([df_cat,df_num],axis=1)
```

In [20]: df

Out[20]:

	Workclass	education	marital-status	occupation	relationship	race	sex	native-country	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	Income
0	5	9	4	0	1	4	1	28	39	77516	13	2174	0	40	0
1	4	9	2	3	0	4	1	28	50	83311	13	0	0	13	0
2	2	11	0	5	1	4	1	28	38	215646	9	0	0	40	0
3	2	1	2	5	0	2	1	28	53	234721	7	0	0	40	0
4	2	9	2	9	5	2	0	4	28	338409	13	0	0	40	0
...
1027	2	11	0	5	3	4	0	28	34	290763	9	0	0	40	0
1028	2	15	2	2	0	4	1	28	36	51100	10	0	0	40	0
1029	2	11	2	13	0	4	1	28	41	227644	9	0	0	50	0
1030	1	9	2	9	5	4	0	28	58	205267	13	0	0	40	1
1031	2	9	2	9	0	1	1	18	53	288020	13	0	0	40	0

1032 rows × 15 columns

In [21]: X = df.iloc[:, :-1]
y = df.iloc[:, -1]

In [22]: X.head()

Out[22]:

	Workclass	education	marital-status	occupation	relationship	race	sex	native-country	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
0	5	9	4	0	1	4	1	28	39	77516	13	2174	0	40
1	4	9	2	3	0	4	1	28	50	83311	13	0	0	13
2	2	11	0	5	1	4	1	28	38	215646	9	0	0	40
3	2	1	2	5	0	2	1	28	53	234721	7	0	0	40
4	2	9	2	9	5	2	0	4	28	338409	13	0	0	40

In [23]: y.head()

Out[23]:

```
0    0
1    0
2    0
3    0
4    0
```

Name: Income, dtype: int64

In [24]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=1)

#Logistic Regression :-

In [49]: log = LogisticRegression(max_iter=120)

In [50]: log.fit(X_train,y_train)

Out[50]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, l1_ratio=None, max_iter=120, multi_class='auto', n_jobs=None, penalty='l2', random_state=None, solver='lbfgs', tol=0.0001, verbose=0, warm_start=False)


```
In [56]: precision_score(y_test,y_pred)
```

```
Out[56]: 0.6785714285714286
```

```
In [57]: recall_score(y_test,y_pred)
```

```
Out[57]: 0.2638888888888889
```

```
In [58]: f1_score(y_test,y_pred)
```

```
Out[58]: 0.38
```

```
In [59]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.81	0.96	0.88	238
1	0.68	0.26	0.38	72
accuracy			0.80	310
macro avg	0.75	0.61	0.63	310
weighted avg	0.78	0.80	0.76	310

```
In [ ]:
```